

Technical Design Document (TDD)

End-to-End ML Classifier Platform with CI/CD

Author: Neha Amin

Audience: CTO, Principal Engineers, Platform & ML Leadership

Status: Work In Progress

Revision: v1.0

1. Executive Summary

This document proposes an upgrade of the existing project into a **production-grade end-to-end machine learning system**:

Data ingestion → Model training → Model registry → REST-based inference → Monitoring → CI/CD automation

The goal is to demonstrate **ML operational readiness**, and **deployment automation**, using **ML infra** while adhering to MLOps best practices.

The system intentionally separates *model development velocity* from *production stability* by decoupling training pipelines from inference deployment.

This system is intentionally designed to be:

- Deterministic and reproducible
 - Auditable and testable
 - CI/CD-driven
 - Cloud-deployable with low operational overhead
-

2. Engineering Objectives

2.1 Engineering Objectives

- Enforce **separation of concerns** across ML lifecycle
 - Apply **DevOps principles to ML workflows (MLOps)**
 - Enable **automated testing and deployment**
 - Ensure low-latency and observable inference
-

3. Scope Definition

In Scope


- Supervised classification model
- Offline training + online inference
- Model versioning
- REST API for predictions
- CI/CD automation
- Basic monitoring and metrics
- Dockerized workloads
- Free-tier cloud deployment

Out of Scope

- Distributed training
 - Real-time feature stores
 - Online learning
 - Advanced monitoring (e.g., drift detection)
 - Enterprise authentication
-

4. High-Level Architecture

4.1 Logical Architecture



4.2 Tech - Stack

Component	Technology
Training	Python + scikit-learn
Registry	Filesystem / MLflow (local)
Inference	FastAPI + Uvicorn
Containerization	Docker
CI/CD	GitHub Actions
Infra	Free-tier Cloud (Render/Fly.io/AWS EC2)
Monitoring	Prometheus-style metrics

5. Detailed Component Design

5.1 Data Ingestion Layer

Purpose:

Ensure clean, validated, reproducible input data.

Responsibilities:

- Load raw datasets
- Validate schema & implies constraints
- Handle missing/invalid labels
- Split data into training and test sets
- Persist processed data artifacts

Output Artifacts:

- `train.csv`
- `test.csv`

Key Engineering Decisions:

- Schema validation prevents silent model degradation
- Deterministic data splits enable reproducibility

5.2 Model Training & Evaluation

Purpose:

Train a deterministic, interpretable classifier with measurable quality.

Responsibilities:

- Feature preprocessing
- Model training
- Metric computation (accuracy, precision, recall)
- Model serialization
- Metadata generation

Model Choice:

- Logistic Regression / Random Forest
(Chosen for interpretability and deterministic behavior)

Artifacts Produced:

- Serialized model (`.joblib`)
 - Metrics summary
 - Dataset hash
-

5.3 Model Registry

Purpose:

Enable versioned and auditable model management.

Design:

- File-based registry
- Semantic versioning (`classifier_v1`, `classifier_v2`)
- Metadata tracking

Metadata Stored:

- Model version
- Training timestamp
- Dataset checksum
- Performance metrics

Rationale:

- Lightweight registry sufficient for small-scale deployment
 - MLflow compatible for future expansion
-

5.4 Inference Service

Purpose:

Expose predictions via a stable REST API.

Framework: FastAPI**Endpoints:****Health Check**

GET /health

Prediction

POST /predict

Response Includes:

- Predicted class
- Confidence score
- Model version

Non-Functional Requirements:

- Latency < 100ms (local)
 - Input validation
 - Stateless service
 - Container-ready
-

5.5 Monitoring & Observability

Purpose:

Provide operational visibility into model behavior.

Metrics Captured:

- Request count
- Inference latency
- Error rate


Endpoint:

GET /metrics

Design Philosophy:

- Lightweight monitoring
- Prometheus-compatible
- Extendable for future drift detection

6. CI/CD Pipeline Design



6.2 CI/CD Stages

Stage 1: Validation

- Python linting
- Unit tests for training & inference
- Schema checks

Stage 2: Training

- Execute training pipeline
- Validate metrics thresholds
- Register new model version

Stage 3: Build

- Build Docker image
- Tag with commit SHA

Stage 4: Deployment

- Push container to cloud
- Restart inference service

Automation Tool: GitHub Actions

7. Containerization Strategy

Docker Principles Applied

- Minimal base image
- Deterministic builds
- Immutable artifacts
- Clear runtime separation

Containers

Container	Responsibility
-----------	----------------

Training	Offline model training
----------	------------------------

Inference	Online prediction
-----------	-------------------

8. Infrastructure (Minimal & Reproducible)

Infrastructure as Code

- Terraform used for provisioning
- Enables repeatable deployments

Deployment Targets

- Render / Fly.io / AWS Free Tier
 - Stateless inference service
-

9. Security & Compliance Considerations

- No secrets committed to source control
 - Environment-based configuration
 - API input validation
 - Minimal exposed surface area
-

10. Testing Strategy

Test Types

- Unit tests (training, inference)
- API contract tests
- Health check validation

CI Enforcement

- Pipeline fails on test failure
 - No manual deployment allowed
-

11. Risks & Mitigations

Risk	Mitigation
Model drift	Versioned retraining
API misuse	Input validation
Cost overrun	Free tier + limits
Reproducibility	Deterministic seeds

12. Success Metrics

- Fully automated CI/CD
 - Sub-100ms inference latency
 - Reproducible model builds
 - Clean rollback via model versioning
 - Deployment without manual steps
-
-

14. Future Enhancements (Roadmap)

- Drift detection
 - Feature store integration
 - Canary deployments
 - Auth-protected inference
 - Centralized model registry
-

15. Conclusion

This upgrade transforms the project from a **demo application** into a **production-ready ML system** aligned with modern **MLOps best practices**.

It demonstrates strong fundamentals in **ML engineering, DevOps, and system design**, making it suitable for real-world deployment and leadership review.

